

- Community-in-the-Loop

Can I Pet That Dawg?

What's safe, normal, or true depends on where you live.
Your AI assumes California.


We put local knowledge back in the loop.



Why this matters: outside English, models work worse, think Western, and protect less

Three measurable failure modes for the people our plugin serves.

1 Quality degrades in every non-English language



All three frontier models peak in English. Arabic, German and Korean trail by 4–7 points — Korean on GPT-5.2 is 13% lower.

2 Cultural bias toward Western norms

−0.70

Correlation between how well GPT reflects a country's real values and its cultural distance from the US. Models echo English-speaking and Protestant-European values — and barely represent Jordan, Pakistan or Libya.

Atari et al. 2023; Tao et al. 2024 (PNAS Nexus)


3 Safety guardrails weaken

+25pp

Rise in unsafe responses when harmful prompts are simply translated into low-resource languages. Safety training is overwhelmingly English.

Welo Data, 2025 (up to 25pp)

⚙️ Two decoupled loops, one knowledge base



Loops are decoupled — feedback is curated async, enrichment reads only verified notes.

Can I pet a dog?

Yes, generally — though it helps to ask the owner first, and let the dog sniff your hand before going in.



Can I pet a dog?

 may vary by culture

Yes, generally — though it helps to ask the owner first, and let the dog sniff your hand before going in.



Where this fits in the AI safety stack

Pre-training data filtering

Before training

RLHF / Constitutional AI

During training

Output guardrails / filters

At inference

← No correction loop post-deployment

Community-in-the-Loop

Post-deployment,
runtime, multilingual

What it does

- ✓ Catches what alignment missed
- ✓ Updates at speed of the world
- ✓ Covers every language spoken
- ✓ Auditable, transparent process

What it doesn't

- ✗ Replace RLHF or guardrails
- ✗ Prevent novel unsafe outputs
- ✗ Work without moderation

Complementary to existing safety layers — adds the first community-driven, post-deployment, multilingual correction loop.

Every safety layer today is lab-controlled
and frozen at training time.

**We add the first community-driven layer
that runs at inference, improves every day,
and speaks every language.**



Community-in-the-Loop



Risks we're thinking about

Mitigated

Cold start — empty knowledge base on day one

Our RAG model adds a configurable cultural prompt to every query — "answer as a person from this culture" — so even before community notes exist, responses are culturally grounded. The community layer improves it over time.

Mitigated

Prompt injection via community notes

Notes are injected as fenced, untrusted reference data — never as instructions. Frontier models are rapidly improving at distinguishing reference context from directives. This is an active area of progress across the industry.

Open

Corpus poisoning — bad data passes moderation

Our consensus gate and reputation system reduce this risk, but a sufficiently patient, coordinated attacker could still get bad notes verified. This is an open problem we haven't fully solved yet.

Future work

Who moderates the moderators?

At scale this needs transparency mechanisms — public audit logs, appeal paths, reviewer diversity requirements. It's a real governance challenge, but it's outside the scope of this hackathon build.

Naming these upfront — every system has risks. Ours are the same ones Wikipedia faced at launch.