

Community-in-the-Loop: Idea Brief

What is the idea?

A tool-agnostic plug-in that allows users to flag regionally/ culturally inaccurate, problematic or dangerous content in LLM responses. Crowd-sourced data gathered is analysed and used to enrich responses, providing a culturally relevant lens on LLM responses.

Community-in-the-Loop will give communities a near-zero-cost way to flag and correct culturally-wrong AI answers, which both improves those answers immediately and builds a living map of where AI fails non-Western users, and that map becomes the evidence that pushes labs and regulators to fix it at source.

It does two things:

- Enriches responses with safe, cultural and local context for the user's country/community before they see them.
- Lets users flag responses that don't fit their culture or context. Those individual flags are aggregated and fed back to improve the enrichment for everyone in that community.

Why is this needed?

- Approx. 52% of content on the web is English, but only 16% are English speakers
 - ([Ethnologue, The Digital Language Divide](#))
- Data shows that performance degrades across different cases
 - MMLU-ProX, a benchmark of 11,829 identical questions translated across 29 languages, found performance drops steadily from high-resource to low-resource languages, with gaps of up to 24.3 percentage points between the best- and worst-served languages. On the strongest model tested (DeepSeek-R1), the best-served languages (French and Indonesian, 81.3%) outscored the worst-served (Yoruba, 57.0%) by exactly that 24.3pp.
- Models systematically homogenize moral diversity
 - (<https://arxiv.org/abs/2507.10073>)
- Models are trained on vast, English-dominated datasets, which flattens global diversity and favor classically western norms
 - A study across 107 countries, benchmarking five widely-used LLMs against national survey data, found their responses cluster around the values of English-speaking and Protestant European countries — even when answering in other languages. ([Tao et al., "Cultural bias and cultural alignment of large language models", PNAS Nexus, 2024](#))
- Safety guardrails - easier to jail-break in non-English language
 - Researchers showed that simply translating an unsafe request into a low-resource language (e.g. Zulu, Scots Gaelic) using free tools bypassed GPT-4's safety filters. Success at eliciting harmful content rose from under 1% in English to around 79% in low-resource languages — on par with state-of-the-art jailbreak methods. The top categories were terrorism, financial

crime, and misinformation. ([Yong, Menghini & Bach, "Low-Resource Languages Jailbreak GPT-4", arXiv 2310.02446](#))

Examples

- **Cultural variation** in perception of dogs.
- **Emergency numbers.** A Western-default model says "call 911." That number does nothing in most of the world — it's 112 across the EU, 999 in the UK, 119 for ambulance in Japan, 000 in Australia, 112/108 in India. In a real emergency, those wasted seconds are the harm.
- **Road and pedestrian norms.** "Pedestrians have the right of way, just use the crosswalk" gets you hit in cities where that norm isn't honored in practice.
- **Carrying common medications across borders.** Adderall and other amphetamine-based ADHD meds are illegal in Japan; codeine and some painkillers are tightly controlled in the UAE and Singapore; even pseudoephedrine cold medicine has caused arrests. "Just pack your prescription" can mean detention.
- **Individualist mental-health framing.** "Set boundaries with your parents," "move out and prioritize yourself," "confront them directly" can backfire badly in strongly collectivist or honor-based family contexts, and "just go to therapy" assumes access and ignores stigma that can affect marriage and employment prospects.

How our solution maps onto the evidence

The problem (evidenced above)

How the plugin addresses it

Performance degrades in low-resource languages (§1)

Enriches prompts/responses with local context the model lacks, at point of use — no retraining needed

Cultural bias toward Western values (§4)

Operationalises "cultural prompting", the fix shown to work for 71–81% of countries — but sources the cultural context from the community instead of guessing

Safety guardrails weaker in other languages (§3)

Community flags surface locally-dangerous or wrong answers (e.g. the rabies example) that the model's safety training never covered

Missing data will never appear online for most cultures (§6)

A human flagging-and-aggregation loop generates exactly the context that the open internet won't

What change will this solution deliver?

- Tailored, nuanced model outputs that reflect cultural values and don't impose western values
- Community led action that challenge current power concentrations

Who is impacted?

- Non-technical, LLM users (primary) - would require link in to AI education stream
- Developers who don't speak English/ English isn't their first language (primary)
- Governments looking to implement community-backed AI governance that is regionally and culturally relevant (secondary)

Is it scalable and how will people access the solution?

- Solution will be in the form of a MCP plug-in
- Must be LLM agnostic, accessible on a mobile and free to use.
- It's scalable and adaptable across different communities.
- Opportunities to complement other initiatives focused on bringing parity to LLM use in different languages

What is out of scope?

Things we're aware of, but aren't directly tackling with this project:

- Higher token costs for languages that aren't as represented in training data sets (lower resourced languages)
 - A peer-reviewed study found tokenisers split the *same* content into very different numbers of tokens across languages — the gap can be as large as 15x between the most and least efficient languages. Because APIs bill per token and models have fixed context limits, this means non-English users pay more and hit limits sooner for identical content. ([Petrov et al., "Language Model Tokenizers Introduce Unfairness Between Languages". NeurIPS 2023, arXiv 2305.15425](#))
- How we protect against bad actors/ malicious use
- Who moderates the moderators

Assumptions

- the communities we most want to serve will actually contribute
- a culture has a consensus "right answer" and that aggregated flags will find it
- MCP becomes the universal connector layer that consumer and mobile AI apps adopt

Why do we think this will work?

- Masakhane (<https://www.masakhane.io/>) - grassroots organisation whose mission is to strengthen and spur NLP research in African languages, for Africans, by Africans.
- Governments are slow to act and policy doesn't keep pace with technology - this tool provides a way for end-users to challenge and flag regionally or culturally irrelevant, problematic or even dangerous responses quickly and easily and will build a corpus of community verified information that governments can then use to support future policy development